

機器學習研究應用

Study for Machine Learning and Its Applications

The Basics of Neural Networks

孫士韋

Shih-Wei Sun

swsun@newmedia.tnua.edu.tw

Outline

- Data representations for neural networks
- The engine of neural networks
- Looking back at our first example

Data Representations for Neural Networks

Data Representations for Neural Networks

- Tensors: container for numbers
 - Google's TensorFlow named after
- Scalars (0D tensors)

- Only **one number**
- Float32 or float 64
- 0 axes (ndim: 0)
- Number of axes: rank

Code:

```
1 import numpy as np
2 x = np.array(12)
3 print(x)
4 print(x.ndim)
```

Result:

12
0

Vectors (1D Tensor)

- Vector: an array of **numbers**
- Exactly **one** axis

Code:

1D tensor: 1 bracket

- The vector

- Has **5 entries**

- 5-dimensional vector

- Only **one axis**

```
6 x = np.array([12, 3, 6, 14, 7])
7 print(x)
8 print(x.ndim)
```

Result:

[12 3 6 14 7]
1D tensor (axis)
5 entries

```
[12  3  6 14  7]
1
```

- **Dimension:** number of **entries**, a **specific axis**

Matrices (2D Tensor)

- Matrix: an array of **vectors**
- **Two** axes

Code:

2D tensor: 2 brackets

The same part of code

```
10 x = np.array([[5, 78, 2, 34, 0],  
11               [6, 79, 3, 35, 1],  
12               [7, 80, 4, 36, 2]])  
13 print(x)  
14 print(x.ndim)
```

1st axis →

↓ 2nd axis

2-D tensor

Result:

```
[[ 5 78  2 34  0]  
 [ 6 79  3 35  1]  
 [ 7 80  4 36  2]]  
2
```

3D Tensor

3-D tensor

- Many arrays: cube
- Three axes

Code:

```
16 x = np.array([[[5, 78, 2, 34, 0],  
17               [6, 79, 3, 35, 1],  
18               [7, 80, 4, 36, 2]],  
19               [[5, 78, 2, 34, 0],  
20               [6, 79, 3, 35, 1],  
21               [7, 80, 4, 36, 2]],  
22               [[5, 78, 2, 34, 0],  
23               [6, 79, 3, 35, 1],  
24               [7, 80, 4, 36, 2]]])  
25 print(x)  
26 print(x.ndim)
```

3D tensor: 3 brackets

3rd axis

2nd axis

1st axis

```
[[[ 5 78  2 34  0]  
 [ 6 79  3 35  1]  
 [ 7 80  4 36  2]]  
 [ 7 80  4 36  2]]
```

Result:

```
[[[ 5 78  2 34  0]  
 [ 6 79  3 35  1]  
 [ 7 80  4 36  2]]  
 [[ 5 78  2 34  0]  
 [ 6 79  3 35  1]  
 [ 7 80  4 36  2]]  
 [[ 5 78  2 34  0]  
 [ 6 79  3 35  1]  
 [ 7 80  4 36  2]]]
```

Key Attributes

- A tensor is defined by:

- Number of axes (rank)

- ndim

- Shape

- How many dimensions the tensor has

- Along each axis

- Data type


- dtype

- float32, float 64, uint8,...

Result:

```
28 print(x.shape) (3, 3, 5)
```

Try to add this line: end of the codes



```
[[[ 5 78  2 34  0]
  [[ 5 78  2 34  0]
  [[ 5 78  2 34  0]
 [ 6 79  3 35  1]
 [ 7 80  4 36  2]]
```

```
[[[ 5 78  2 34  0]
  [[ 5 78  2 34  0]
  [[ 5 78  2 34  0]
 [ 6 79  3 35  1]
 [ 7 80  4 36  2]]
```

3x3x5

Shape: (3,3,5)

3 slices of
3x5 arrays

Like a toast: different slices

Sesame, peanuts, raisin

An Example of Key Attributes

- Load the data from MNIST
- Display the
 - number of axes: the tensor of `train_image`, `ndim`
 - shape of the tensor
 - datatype

Result:

Using TensorFlow backend.

3


3D tensor

(60000, 28, 28)

uint8

60,000 matrices
of 28x28 integers
(8-bit)

60,000
slices



Each matrix:
grayscale image, 0~255 28x28

Code:

```
1 from tensorflow.keras.datasets import mnist
2 (train_images, train_labels),
  (test_images, test_labels) =
  mnist.load_data()
```

```
3
4 print(train_images.ndim)
5 print(train_images.shape)
6 print(train_images.dtype)
```

Displaying the fourth digit

- Display the 4th digit in this 3D tensor

- Using the library **Matplotlib**

– Standard scientific Python suite

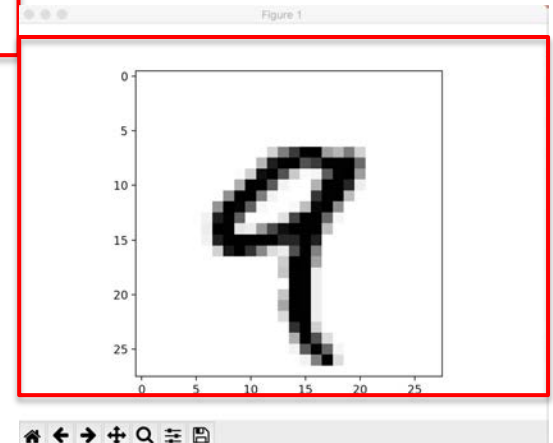
60,000 matrices
of 28x28 integers
(8-bit)

Code:

(60000, 28, 28)

```
8 import matplotlib.pyplot as plt
9 digit = train_images[4]
10 plt.imshow(digit, cmap=plt.cm.binary)
11 plt.show()
```

Result:








Practice 2

- Based on the above code
- Use a for in range loop
 - Hint:

```
for x in range(2, 6):  
    print(x)
```
- Display the 5th to the 10th digits
 - of the MNIST dataset

```
Last login:  
Sunde-iMac:~  
2  
3  
4  
5  
Sunde-iMac:~
```



The gears of neural networks

Manipulating Tensors

- Select a specific digit
 - Alongside the first axis: `train_images[i]`
- Select specific elements in a tensor
 - Tensor slicing
- Select digits #10 to #100 (#100 not included)

Code:

```
13 my_slice = train_images[10:100]
14 print(my_slice.shape)
```

Result:

```
Using TensorFlow backend.
(90, 28, 28)
```

90 slices of 28x28 matrices

Examples of Tensors Slicing

- Equivalent to the previous example

Code:


```
13 my_slice = train_images[10:100]
14 print(my_slice.shape)
15
16 my_slice = train_images[10:100, :, :]
17 print(my_slice.shape)
18
19 my_slice = train_images[10:100, 0:28, 0:28]
20 my_slice.shape
```

Result:

(90, 28, 28)
(90, 28, 28)
(90, 28, 28)

The notation of data batches

- The first axis (axis 0, index from 0)
 - Sample axis => batch axis
 - in MNIST, sample: images of digits
- Deep learning models
 - Don't process an entire dataset at once
(a small group of data)
- Break the data into small batches



Keras, deep-learning libraries

```
model.fit(train_images, train_labels,  
          epochs=5, batch_size=128)
```

- Batch example: batch size of 128

Code:

```
22 batch = train_images[:128]  
23 print(batch.shape)  
24 batch = train_images[128:256]  
25 print(batch.shape)
```

Result:

First batch (128, 28, 28)
(128, 28, 28)

next batch

More batches...

Full code

```
1 from tensorflow.keras.datasets import mnist
2 (train_images, train_labels),
   (test_images, test_labels) =
   mnist.load_data()
3
4 print(train_images.ndim)
5 print(train_images.shape)
6 print(train_images.dtype)
7
8 import matplotlib.pyplot as plt
9 digit = train_images[4]
10 plt.imshow(digit, cmap=plt.cm.binary)
11 plt.show()
12
13 my_slice = train_images[10:100]
14 print(my_slice.shape)
15
16 my_slice = train_images[10:100, :, :]
17 print(my_slice.shape)
18
19 my_slice = train_images[10:100, 0:28, 0:28]
20 my_slice.shape
21
22 batch = train_images[:128]
23 print(batch.shape)
24 batch = train_images[128:256]
25 print(batch.shape)
```


Real-world examples of data tensors

- Vector data
 - 2D tensor (samples, features)
- Timeseries or sequence data
 - 3D tensor (samples, timesteps, feature)
- Images
 - 4D tensors (samples, channels, height, width)
- Video
 - 5D tensors (samples, frames, channels, height, width)

Figure 2.3 A rank-3 timeseries data tensor





Figure 2.4 A rank-4 image data tensor (channels-first convention)

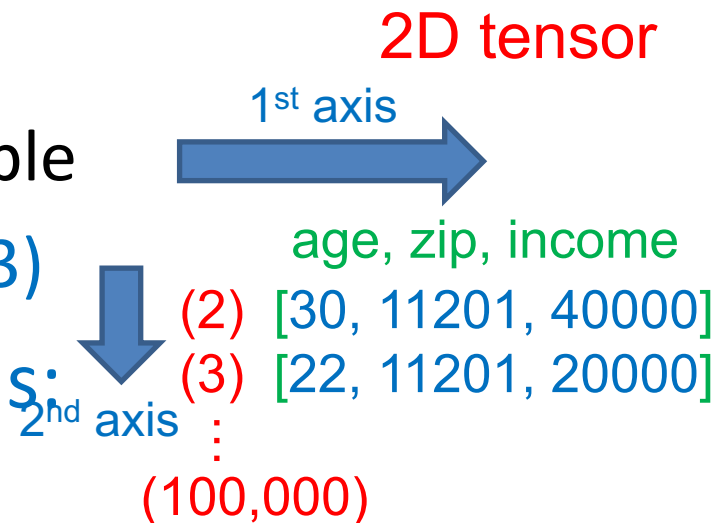
4D tensor



Examples of a 2D tensor

- Actuarial dataset of people
 - Person's age
 - ZIP code
 - Income
- Each person characterized as a vector of
 - 3 values (feature)
 - Entire dataset of 100,000 people
 - 2D tensor of shape: (100000, 3)

- Example of #2 and #3 persons:



Timeseries data or sequence data, 3D Tensor (audio, sensor data)

- Stock price dataset

- For every minute

- Current price of the stock
 - Highest price in the past minute
 - Lowest price in the past minute

- Entire day of trading: 390 minutes

- 2D tensor of shape (390, 3)

- 250 trading days:

- (250, 390, 3)

- Each sample: one day's worth of data





Image data

- Images
 - Height
 - Width
 - Color depth
 - Gray scale image have only **single** color channel
- Example of an **256x256** color image
 - with a batch size of **128**
 - (**128**, **256**, **256**, **3**): (**samples**, **height**, **width**, **color_depth**)
 - Channels-last in **Tensorflow**
 - **Channel-first in Theano**: (**samples**, **color_depth**, **height**, **width**)
 - Both conventions supported: **Keras**



Tensor reshaping

- Preprocessing the digits data
 - Before feeding it into the neural network

```
# look at the training data
```

Using TensorFlow backend.

```
print(train_images.shape)
```

(60000, 28, 28)

```
train_images = train_images.reshape((60000, 28 * 28))
```

- Reshaping: rearranging its rows and columns
 - To match a target shape: (28, 28) 2D to (28*28) 1D

- Same total number
 - As the initial tensor
- ```
print(train_images.shape)
```
- (60000, 784)

# Example of Tensor Reshaping

- A simple example

Code:

```
1 import numpy as np
2
3 x = np.array([[0., 1.],
4 [2., 3.],
5 [4., 5.]])
6 print(x.shape)
7 print(x)
8
9 x = x.reshape((6, 1))
10 print(x.shape)
11 print(x)
12
13 x = x.reshape((2, 3))
14 print(x.shape)
15 print(x)
```

Result:

```
(3, 2)
[[0. 1.]
 [2. 3.]
 [4. 5.]
```

```
(6, 1)
[[0.]
 [1.]
 [2.]
 [3.]
 [4.]
 [5.]
```

```
(2, 3)
[[0. 1. 2.]
 [3. 4. 5.]
```

# Practice 3

- Run the above code
  - For tensor reshaping

Result:

```
(3, 2)
[[0. 1.]
 [2. 3.]
 [4. 5.]]
```

```
(6, 1)
[[0.]
 [1.]
 [2.]
 [3.]
 [4.]
 [5.]]
```

```
(2, 3)
[[0. 1. 2.]
 [3. 4. 5.]]
```

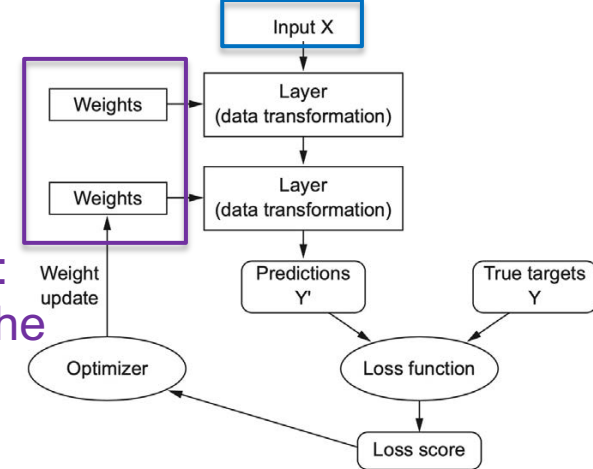
The engine of neural networks



# Tensor Operations

neural  
network

W (weights):  
learned by the  
network



- In our Initial example:

- `keras.layers.Dense(512, activation='relu')`


- The layer: function

- Input: 2D tensor, **output**: 2D tensor

- The function is like this:

relu: activation function

- **output** = `relu(dot(W, input) + b)`



- Three tensor operations

W: weights / (kernel)

b: trainable parameters / (bias)

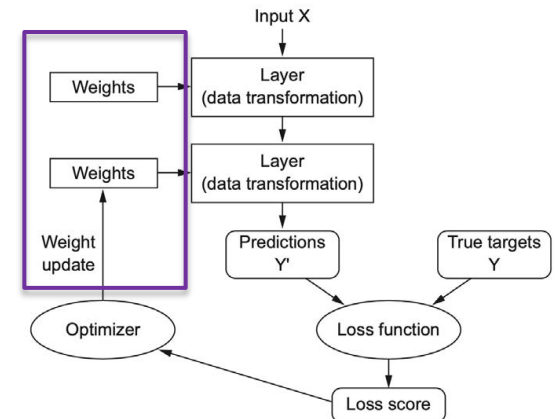
- a dot product (**dot**) between the **input**, **W (weight)**

- addition (+): resulting 2D tensor and a vector **b (bias)**

- **relu** operation: `relu(x)` is `max(x, 0)`

# Weight matrix $W$

- function:  $\text{output} = \text{relu}(\text{dot}(W, \text{input}) + b)$
- Weight matrix  $W$  :
  - filled with small random values (random init.)
  - No reason to expect the function
    - Useful representations, just the starting point!
- Gradual adjustment: training in machine learning!
- Gradual adjust the weights  $W$ 
  - Based on a feedback signal

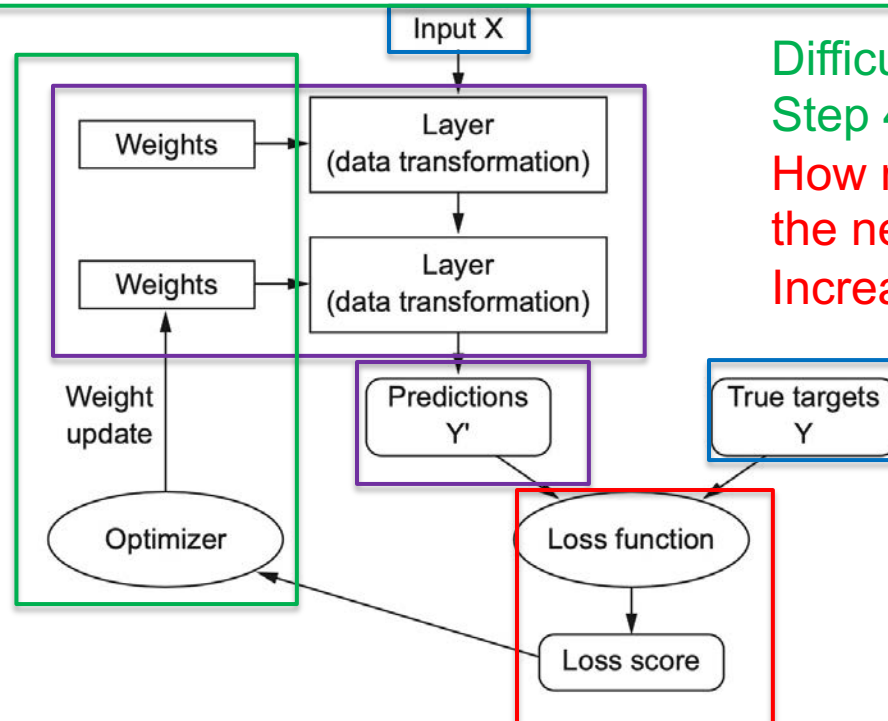


# Training loop

- 1 Draw a batch of training samples  $x$  and corresponding targets  $y$ .
- 2 Run the network on  $x$  (a step called the *forward pass*) to obtain predictions  $y_{\text{pred}}$ .
- 3 Compute the loss of the network on the batch, a measure of the mismatch between  $y_{\text{pred}}$  and  $y$ .
- 4 Update all weights of the network in a way that slightly reduces the loss on this batch.

- Finally:

- End up a network
  - With low loss
- Low mismatch
  - $y_{\text{pred}}$  and  $y$
- Network learned

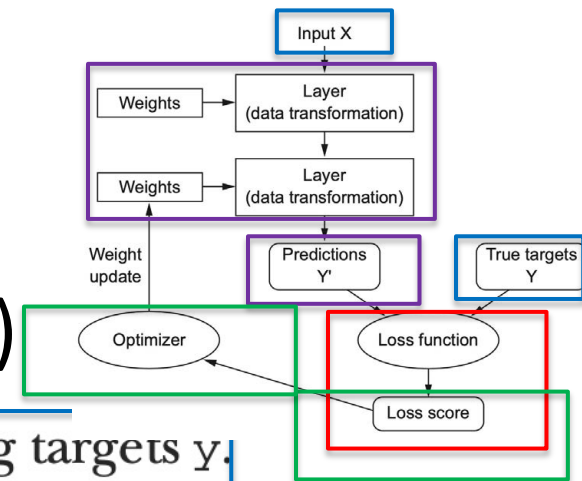


Difficult part:  
Step 4

How much to update  
the network's weight?  
Increase or decrease?

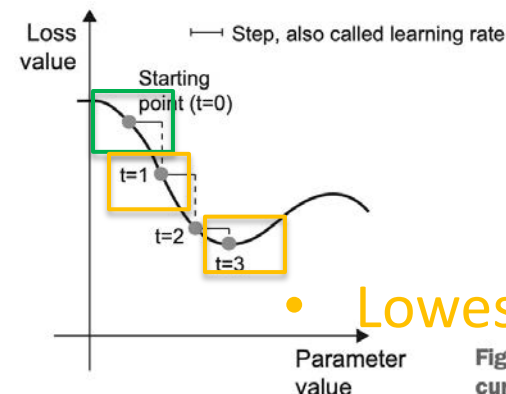
# Optimization

- Stochastic gradient decent (SGD)



- 1 Draw a batch of training samples  $x$  and corresponding targets  $y$ .
- 2 Run the network on  $x$  to obtain predictions  $y_{\text{pred}}$ .
- 3 Compute the loss of the network on the batch, a measure of the mismatch between  $y_{\text{pred}}$  and  $y$ .
- 4 Compute the gradient of the loss with regard to the network's parameters (a *backward pass*).
- 5 Move the parameters a little in the opposite direction from the gradient—for example  $W -= \text{step} * \text{gradient}$ —thus reducing the loss on the batch a bit.


- Move the parameters
  - until Gradient = 0
  - (Derivative = 0)



- Lowest value: loss

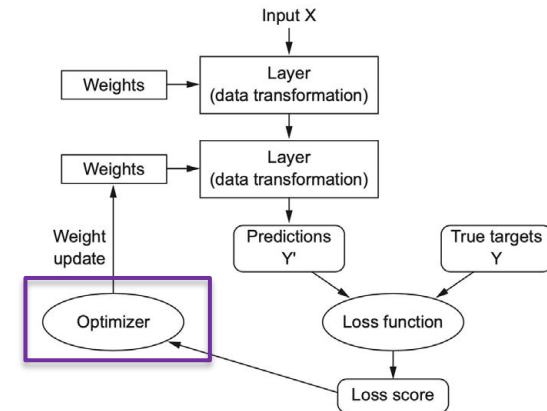
# Local minimum vs. a global minimum

- Stuck in a local minimum
- Need a momentum
  - Moving a ball by a velocity



Rmsprop (optimizer):  
Root-mean-square propagation algorithm  
SGD + momentum

To: Get the smallest LOSS



**Figure 2.13** A local minimum and a global minimum

Looking back at our first example

# Looking back at our first example

- Input data

```
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype("float32") / 255
test_images = test_images.reshape((10000, 28 * 28))
test_images = test_images.astype("float32") / 255
```

- Network

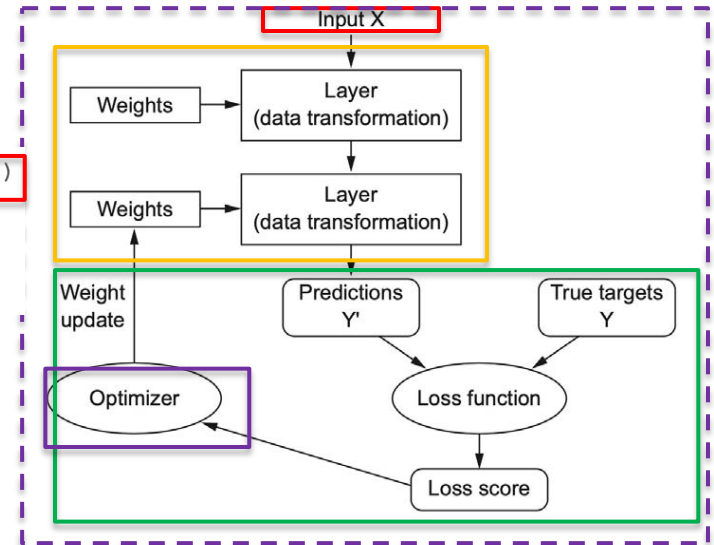
```
model = keras.Sequential([
 layers.Dense(512, activation="relu"),
 layers.Dense(10, activation="softmax")
)
```

- Network-compilation

```
model.compile(optimizer="rmsprop",
 loss="sparse_categorical_crossentropy",
 metrics=["accuracy"])
```

- Training

```
model.fit(train_images, train_labels, epochs=5, batch_size=128)
```



Loss and optimizer: you need to define  
- Before feeding data into a network

**loss:** quantity attempt to minimize

**optimizer:** gradient of the loss,  
Uses to update parameters

Epoch: a full training loop  
Epochs: number of training loops

How many times for staying in a training loop?

# Textbook Reading

- *Deep Learning with Python*
  - Ch 2, the mathematical building blocks of neural networks
    - p.31 - p. 62