

# 機器學習研究應用

## Study for Machine Learning and Its Applications

### Fundamentals of Machine Learning

孫士韋

Shih-Wei Sun

[swsun@newmedia.tnua.edu.tw](mailto:swsun@newmedia.tnua.edu.tw)

# Outline

- Four branches of machine learning
- Evaluating machine-learning models
- Data preprocessing, feature engineering, and feature learning
- Overfitting and underfitting
- The universal workflow of machine learning

# 4 Branches of machine learning (1/2)

- **Supervised learning** (goal of this course)

- Binary classification: **IMDB: like, dislike**
- multiclass classification: **Reuters news, 46 classes**
- Scalar regression: **Boston housing price**
- **Learn the training input / training targets (label, annotation)**

- **Examples:**

- Optical character recognition, Speech recognition
- Image classification, Language translation
- **Sequence generation**: given a image, predict a caption
- **Object detection**: given a picture, draw a bounding box
- **Image segmentation**: given a picture, draw a mask

Car,  
[Sun et al., JVCI, 13]



Segmentation



chair




# 4 Branches of machine learning (2/2)

- Supervised learning (major part)
- Unsupervised learning Image compression
  - Data visualization, **compression**, denoising
    - Dimensionality reduction, clustering (math field)
- Self-supervised learning
  - Supervised learning: **without human-annotated labels**
  - **labels**: **generated from the input data, autoencoder**
    - Predict the next frame in a video / given past frames
    - Predict the next word in a text / given previous words
- Reinforcement learning
  - **Google DeepMind, Alpha Go**, agent
    - Receives information about its environment
    - Learns to choose actions, maximize some reward



Self-driving car



robot

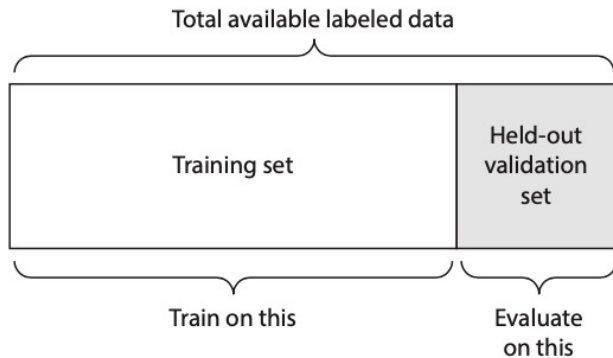


# Evaluating machine-learning models

- Generalize the model
  - Never-before-seen data
- Training, validation, test sets
  - Splitting the available data into 3 sets
  - Avoid information leaks
    - Training, validation data: similar manner
  - Test dataset: never-before-seen
    - New data, ex: interactive hand gesture, talk to Siri

# Validations

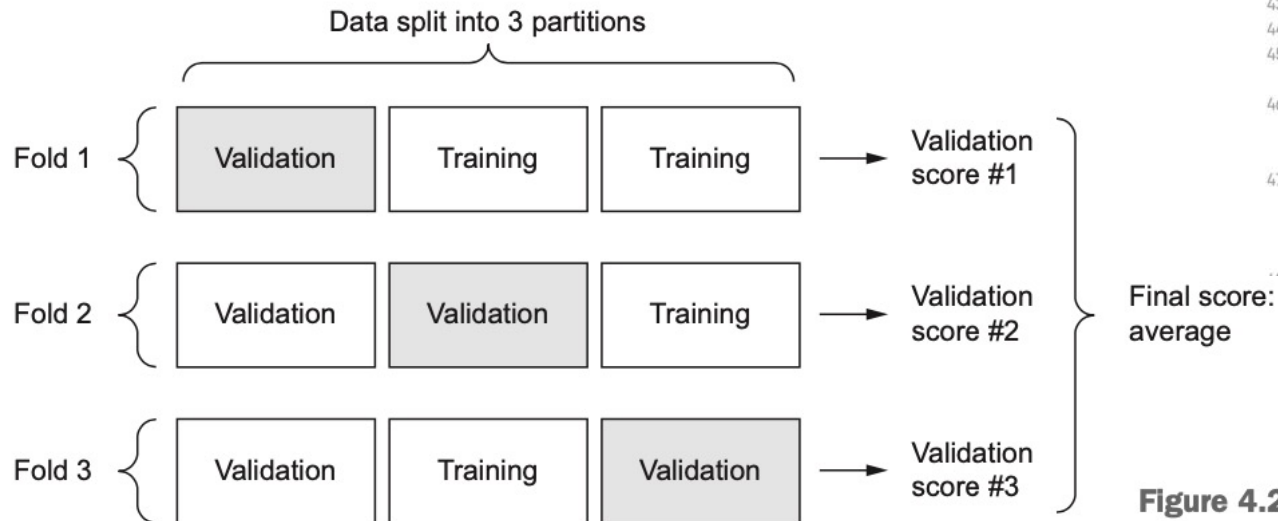
- Simple hold-out validation: early examples



**Figure 4.1** Simple hold-out validation split

```
6 my_slice = train_images[10:100, :, :]
7 print(my_slice.shape)
8
9 my_slice = train_images[10:100, 0:28, 0:28]
10 print(my_slice.shape)
```

- K-fold validation: previous example (last week)



```
43 for i in range(k):
44     print('processing fold #%d' % i)
45     val_data = train_data[i * num_val_samples:
46                           (i + 1) * num_val_samples]
46     val_targets = train_targets[i *
47                                num_val_samples: (i + 1) *
48                                num_val_samples]
47     partial_train_data =
48         np.concatenate([train_data[i *
49                                   num_val_samples: (i + 1) *
50                                   num_val_samples:],
51                           train_data[(i + 1) *
52                                   num_val_samples:]),
53                           axis=0)
```

**Figure 4.2** Three-fold validation

# Things to keep in mind

- Data representativeness
  - In a digit recognition example
    - First 80%: training set
    - Remaining 20%: testing set
    - Model: 0-7
    - Test for 8,9: ridiculous!
  - Random shuffle the samples!
- Arrow of time
  - Predict the future / given the past
  - Should not shuffle your data
- Redundancy: data points appear twice, avoid

# Data preprocessing for neural networks (1/2)

- Vectorization
  - Must be **Tensors** of **float point data**
    - Sound, image: **float32 data**
    - Text: sequence of words: **one-hot encoding**
- Value normalization
  - Image data: 0 – 255 range (gray scale values)
    - **Divide by 255**: values in **0 – 1** range
  - House price:
    - **Standard deviation of 1** and **mean of 0**

```
x -= x.mean(axis=0)  
x /= x.std(axis=0)
```



Assuming x is a 2D data matrix  
of shape (samples, features)



# Data preprocessing for neural networks (2/2)

- Handling missing values

- Sometimes missing some data

- Set it as 0

- 0: means missing data, start ignoring the value

- Copy some training samples

- Feature engineering

- Input an image

- Output the time of a day

- Input

- **Raw pixels**: a bad way

- **Feature values**: better

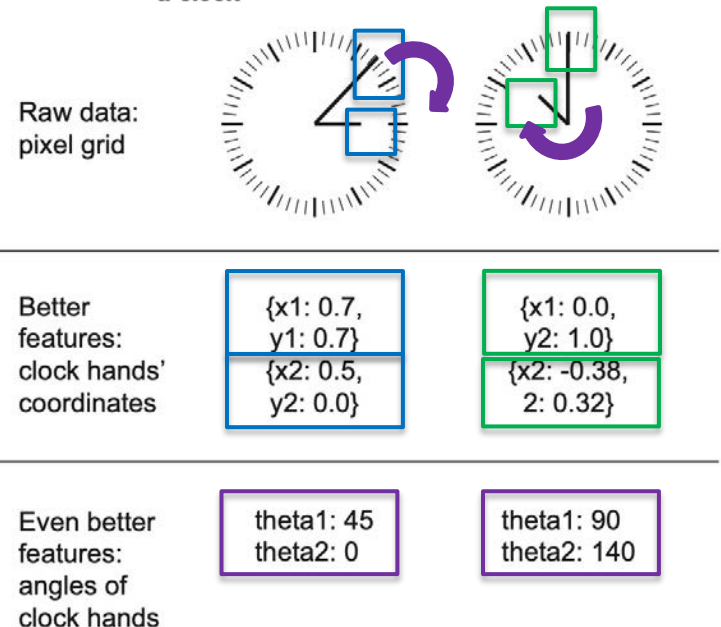
Feature engineering:

Critical role!!

-Make to problem solved or not

-Far less data

Figure 4.3 Feature engineering for reading the time on a clock



# In the previous examples

- **Overfit** happens in every
  - Machine learning problem
- Optimization The thing we can control

- **Adjusting a model**


- To get the best performance

- Generalization
  - How well of the performance
  - Target: get good generalization

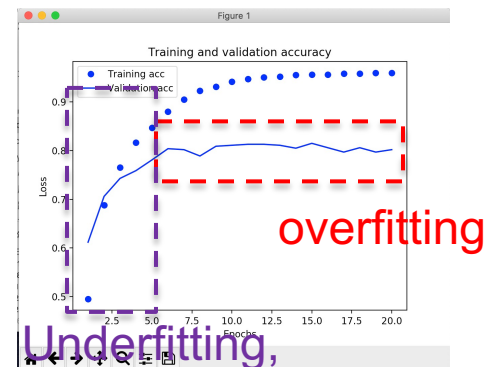
- The **best** solution for **overfitting**
  - **Get more training data!**

Get naturally generalized. But...We don't have it often.

IMDB binary classification,  
accuracy



Reuters Newswire,  
multiclass classification,  
accuracy



Underfitting,  
more epochs,  
more learned parameters,  
more accuracy

# Skills to Fight Overfitting

- The **overfitting problem**:
  - The model **memorized too many things**
    - From the training data
    - Try to **remove** the **too many memorized things**
- To make the model
  - **More generalized**
- 3 methods:
  - Reducing the network's size
  - Adding weight regularization
  - Adding dropout

# Method 1:

## Reducing the network's size

- The simplest way to prevent overfitting
  - reduce number of learnable parameters in the model: 16 -> 4
  - Reduce the memorization capacity

### Listing 4.3 Original model

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

No magical formula to determine: right number

Try to start from a small number

### Listing 4.4 Version of the model with lower capacity

```
model = models.Sequential()
model.add(layers.Dense(4, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(4, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

# Starting from an IMDB Example (1/2)

- Load the data / label
- Transform texts to integers

```
1 from tensorflow.keras.datasets import imdb
2 import numpy as np
3
4 (train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=10000)
```

Load the data, label

```
6 def vectorize_sequences(sequences, dimension=10000):
7     # Create an all-zero matrix of shape (len(sequences), dimension)
8     results = np.zeros((len(sequences), dimension))
9     for i, sequence in enumerate(sequences):
10         results[i, sequence] = 1. # set specific indices of results[i] to 1s
11     return results
12
```

Transform texts to integers

# Starting from an IMDB Example

## (2/2)

- Preparing the data

```
13 # Our vectorized training data
14 x_train = vectorize_sequences(train_data)
15 # Our vectorized test data
16 x_test = vectorize_sequences(test_data)
17 # Our vectorized labels
18 y_train = np.asarray(train_labels).astype('float32')
19 y_test = np.asarray(test_labels).astype('float32')
20
```

Call the function to  
Transform texts to integers

labels  
to  
integers

# Setting up for the Neural Network

- Step 1: add
- Step 2: compile
- Step 3: fit

```
21 from tensorflow.keras import models
22 from tensorflow.keras import layers
```

importing models, layers

```
23
24 #reducing the network's size
```

Step1:  
add  
layers

Step 2:  
Compile,  
Setup the configuration

```
25 original_model = models.Sequential()
```

```
26 original_model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
```

```
27 original_model.add(layers.Dense(16, activation='relu'))
```

```
28 original_model.add(layers.Dense(1, activation='sigmoid'))
```

```
29 original_model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
```

```
30
```

```
31 original_hist = original_model.fit(x_train, y_train, epochs=20, batch_size=512,
    validation_data=(x_test, y_test))
```

Step 3: fit the model

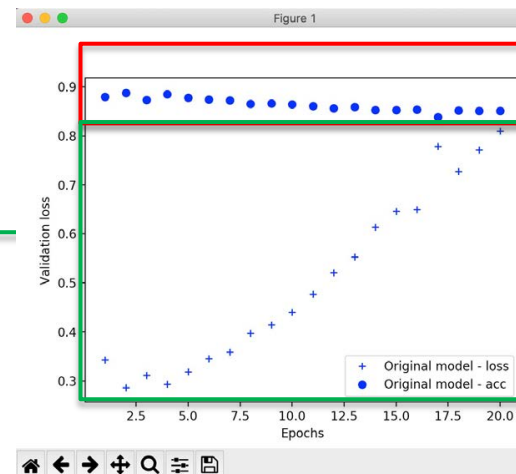
# Plot the Overfitting Curve

- Plot the results

```
--
33 epochs = range(1, 21)
34 original_val_loss = original_hist.history['val_loss']
35
36 original_val_acc = original_hist.history['val_acc']
37
38 import matplotlib.pyplot as plt
39
40 # b+ is for "blue cross"
41 plt.plot(epochs, original_val_loss, 'b+', label='Original model - loss')
42
43 plt.plot(epochs, original_val_acc, 'bo', label='Original model - acc')
44
45 plt.xlabel('Epochs')
46 plt.ylabel('Validation loss')
47 plt.legend()
48 plt.show()
49
```

Assign  
Validation  
loss, acc

Plot  
the  
results



Accuracy  
unstable,  
decreased

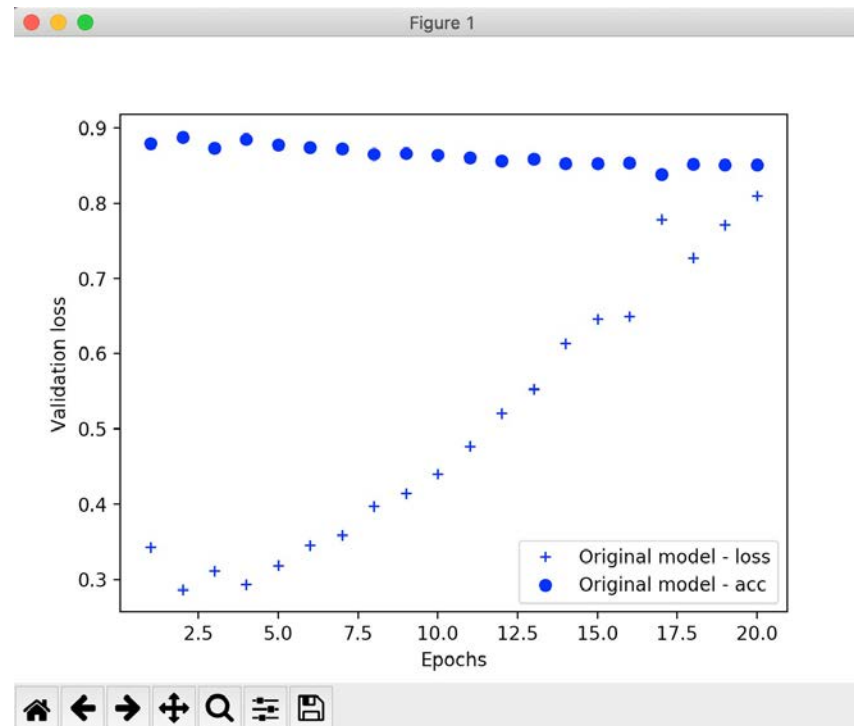
Loss increased  
-not decrease  
(good way),

```
Epoch 3/20
25000/25000 [=====] - 2s 77us/step - loss: 0.2619 - acc
: 0.9214 - val_loss: 0.3364 - val_acc: 0.8842
Epoch 4/20
25000/25000 [=====] - 3s 112us/step - loss: 0.2407 - ac
c: 0.9307 - val_loss: 0.3681 - val_acc: 0.8719
Epoch 5/20
25000/25000 [=====] - 3s 114us/step - loss: 0.2286 - ac
c: 0.9364 - val_loss: 0.3667 - val_acc: 0.8745
Epoch 6/20
25000/25000 [=====] - 2s 88us/step - loss: 0.2212 - acc
: 0.9392 - val_loss: 0.3767 - val_acc: 0.8718
Epoch 7/20
22016/25000 [=====>....] - ETA: 0s - loss: 0.2164 - acc: 0.9
422
```



# Practice 1

- Plot the overfitting curve



# Method 1:

## Reducing the network's size

- Network size: 16 -> 4

Step1: add layers,  
Smaller network's size  
16 -> 4

Step 2: Compile  
Setup the  
configuration

```
50 #method 1: reducing the network's size
51 smaller_model = models.Sequential()
52 smaller_model.add(layers.Dense(4, activation='relu', input_shape=(10000,)))
53 smaller_model.add(layers.Dense(4, activation='relu'))
54 smaller_model.add(layers.Dense(1, activation='sigmoid'))
55 smaller_model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
56
57 smaller_model_hist = smaller_model.fit(x_train, y_train, epochs=20, batch_size=512,
58                                       validation_data=(x_test, y_test))
59 smaller_model_val_loss = smaller_model_hist.history['val_loss']
```

Step 3: Fit the model


Assigning the loss results

# Plot the network reducing results

- Plot the results

```
60
61 # b+ is for "blue cross"
62 plt.plot(epochs, original_val_loss, 'b+', label='Original model')
63 # "bo" is for "blue dot"
64 plt.plot(epochs, smaller_model_val_loss, 'bo', label='Smaller model')
65 plt.xlabel('Epochs')
66 plt.ylabel('Validation loss')
67 plt.legend()
68
69 plt.show()
70
```

```
Epoch 3/20
25000/25000 [=====] - 2s 77us/step - loss: 0.2619 - acc
: 0.9214 - val_loss: 0.3364 - val_acc: 0.8842
Epoch 4/20
25000/25000 [=====] - 3s 112us/step - loss: 0.2407 - ac
c: 0.9307 - val_loss: 0.3681 - val_acc: 0.8719
Epoch 5/20
25000/25000 [=====] - 3s 114us/step - loss: 0.2286 - ac
c: 0.9364 - val_loss: 0.3667 - val_acc: 0.8745
Epoch 6/20
25000/25000 [=====] - 2s 88us/step - loss: 0.2212 - acc
: 0.9392 - val_loss: 0.3767 - val_acc: 0.8718
Epoch 7/20
22016/25000 [=====>....] - ETA: 0s - loss: 0.2164 - acc: 0.9
422
```

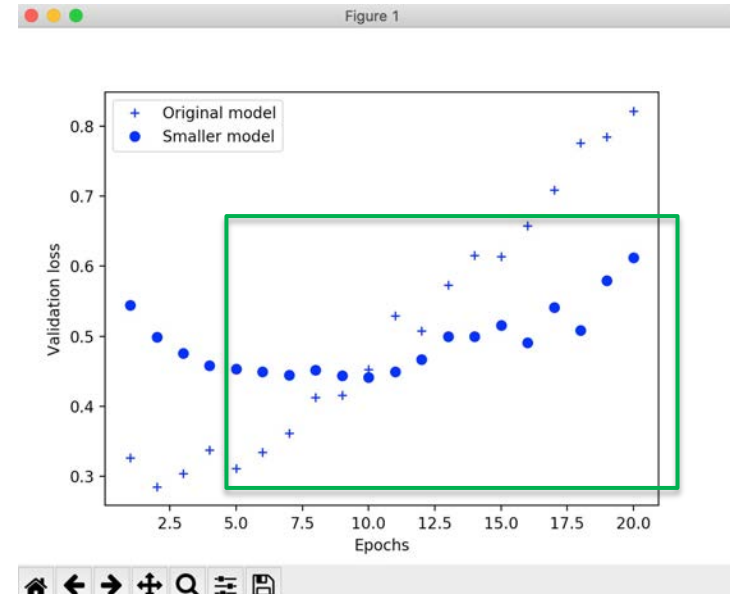



# Practice 2

- Try other numbers

## Listing 4.4 Version of the model with lower capacity

```
model = models.Sequential()  
model.add(layers.Dense(4, activation='relu', input_shape=(10000,)))  
model.add(layers.Dense(4, activation='relu'))  
model.add(layers.Dense(1, activation='sigmoid'))
```



# Method 2:

## Adding weight regularization

- A Simple model
  - Less likely to overfit than complex ones
  - Distribution of parameter values
    - Less entropy (a model with fewer parameters)
- Put constraints
  - Forcing weights: small values
  - Make the distribution of weight values
    - More regular – weight regularization
- L1 regularization: L1 norm, absolute value
- L2 regularization: L2 norm, square value

Adding the loss function, Cost added

of the weight coefficients

# Method 2:

## Adding weight regularization

- Add weight **regularizer**

```
71 # method 2: adding weight regularization
72 from tensorflow.keras import regularizers
73
74 l2_model = models.Sequential()
75 l2_model.add(layers.Dense(16, kernel_regularizer=regularizers.l2(0.001),
76                             activation='relu', input_shape=(10000,)))
76 l2_model.add(layers.Dense(16, kernel_regularizer=regularizers.l2(0.001),
77                             activation='relu'))
77 l2_model.add(layers.Dense(1, activation='sigmoid'))
78
79 l2_model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
80 l2_model_hist = l2_model.fit(x_train, y_train, epochs=20, batch_size=512,
81                             validation_data=(x_test, y_test))
82
83 l2_model_val_loss = l2_model_hist.history['val_loss']
```

Step 1: add layers,  
With regularizers, l2(0.001)

Step 2: Compile  
Setup the  
configuration

Step 3:  
Fit the model


Assigning the loss results

# Plot the regularization results

- Plot the results


```
84 plt.plot(epochs, original_val_loss, 'b+', label='Original model')
85 plt.plot(epochs, l2_model_val_loss, 'bo', label='L2-regularized model')
86 plt.xlabel('Epochs')
87 plt.ylabel('Validation loss')
88 plt.legend()
89
90 plt.show()
91
```

```
Epoch 3/20
25000/25000 [=====] - 2s 77us/step - loss: 0.2619 - acc
: 0.9214 - val_loss: 0.3364 - val_acc: 0.8842
Epoch 4/20
25000/25000 [=====] - 3s 112us/step - loss: 0.2407 - ac
c: 0.9307 - val_loss: 0.3681 - val_acc: 0.8719
Epoch 5/20
25000/25000 [=====] - 3s 114us/step - loss: 0.2286 - ac
c: 0.9364 - val_loss: 0.3667 - val_acc: 0.8745
Epoch 6/20
25000/25000 [=====] - 2s 88us/step - loss: 0.2212 - acc
: 0.9392 - val_loss: 0.3767 - val_acc: 0.8718
Epoch 7/20
22016/25000 [=====>....] - ETA: 0s - loss: 0.2164 - acc: 0.9
422
```



# Practice 3

- Plot the network reducing results





# Method 3:

## Adding dropout

- Dropout
  - One of the most effective way,
    - commonly used
  - Regularization techniques
  - [Geoff Hinton et al, 14], U of Toronto
  - Applied to a layer (setting to zero)

Neural network matrix

0.3	0.2	1.5	0.0
0.6	0.1	0.0	0.3
0.2	1.9	0.3	1.2
0.7	0.5	1.0	0.0

50% dropout

0.0	0.2	1.5	0.0
0.6	0.1	0.0	0.3
0.0	1.9	0.3	0.0
0.7	0.0	0.0	0.0

\* 2

Strange and arbitrary,  
But really work!

Introduce some noise

**Figure 4.8** Dropout applied to an activation matrix at training time, with rescaling happening during training. At test time, the activation matrix is unchanged.

# Method 3: adding dropout

- Dropout

Step1: add layers,  
With dropout

```
92 #method 3: adding dropout
```

```
93 dpt_model = models.Sequential()  
94 dpt_model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))  
95 dpt_model.add(layers.Dropout(0.5))  
96 dpt_model.add(layers.Dense(16, activation='relu'))  
97 dpt_model.add(layers.Dropout(0.5))  
98 dpt_model.add(layers.Dense(1, activation='sigmoid'))
```

```
99  
100 dpt_model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
```

Step 2: Compile, Setup the configuration

```
101  
102 dpt_model_hist = dpt_model.fit(x_train, y_train, epochs=20, batch_size=512,  
    validation_data=(x_test, y_test))
```

```
103  
104 dpt_model_val_loss = dpt_model_hist.history['val_loss']
```

Step 3:  
Fit the model

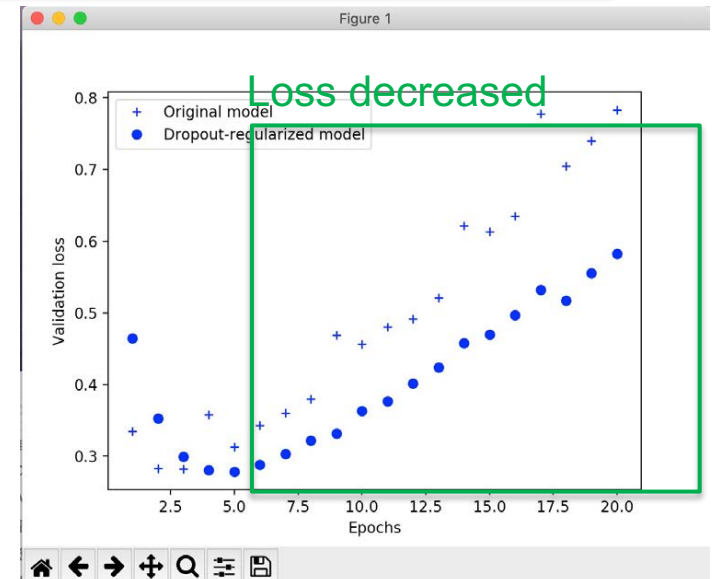
Assigning the loss results

# Plot the network reducing results

- Plot the results


```
105 plt.plot(epochs, original_val_loss, 'b+', label='Original model')
106 plt.plot(epochs, dpt_model_val_loss, 'bo', label='Dropout-regularized model')
107 plt.xlabel('Epochs')
108 plt.ylabel('Validation loss')
109 plt.legend()
110
111
112 plt.show()
```

```
Epoch 3/20
25000/25000 [=====] - 2s 77us/step - loss: 0.2619 - acc
: 0.9214 - val_loss: 0.3364 - val_acc: 0.8842
Epoch 4/20
25000/25000 [=====] - 3s 112us/step - loss: 0.2407 - ac
c: 0.9307 - val_loss: 0.3681 - val_acc: 0.8719
Epoch 5/20
25000/25000 [=====] - 3s 114us/step - loss: 0.2286 - ac
c: 0.9364 - val_loss: 0.3667 - val_acc: 0.8745
Epoch 6/20
25000/25000 [=====] - 2s 88us/step - loss: 0.2212 - acc
: 0.9392 - val_loss: 0.3767 - val_acc: 0.8718
Epoch 7/20
22016/25000 [=====>....] - ETA: 0s - loss: 0.2164 - acc: 0.9
422
```



# Practice 4

- Plot the dropout results



# Full code (1/5)

```
1 from tensorflow.keras.datasets import imdb
2 import numpy as np
3
4 (train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=10000)
5
6 def vectorize_sequences(sequences, dimension=10000):
7     # Create an all-zero matrix of shape (len(sequences), dimension)
8     results = np.zeros((len(sequences), dimension))
9     for i, sequence in enumerate(sequences):
10         results[i, sequence] = 1. # set specific indices of results[i] to 1s
11     return results
12
13 # Our vectorized training data
14 x_train = vectorize_sequences(train_data)
15 # Our vectorized test data
16 x_test = vectorize_sequences(test_data)
17 # Our vectorized labels
18 y_train = np.asarray(train_labels).astype('float32')
19 y_test = np.asarray(test_labels).astype('float32')
20
21 from tensorflow.keras import models
22 from tensorflow.keras import layers
23
```

# Full code (2/5)

```
24 #reducing the network's size
25 original_model = models.Sequential()
26 original_model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
27 original_model.add(layers.Dense(16, activation='relu'))
28 original_model.add(layers.Dense(1, activation='sigmoid'))
29 original_model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
30
31 original_hist = original_model.fit(x_train, y_train, epochs=20, batch_size=512,
    validation_data=(x_test, y_test))
32
33 epochs = range(1, 21)
34 original_val_loss = original_hist.history['val_loss']
35
36 original_val_acc = original_hist.history['val_acc']
37
38 import matplotlib.pyplot as plt
39
40 # b+ is for "blue cross"
41 plt.plot(epochs, original_val_loss, 'b+', label='Original model - loss')
42
43 plt.plot(epochs, original_val_acc, 'bo', label='Original model - acc')
44
45 plt.xlabel('Epochs')
46 plt.ylabel('Validation loss')
47 plt.legend()
48 plt.show()
```

# Full code (3/5)

```
50 #method 1: reducing the network's size
51 smaller_model = models.Sequential()
52 smaller_model.add(layers.Dense(4, activation='relu', input_shape=(10000,)))
53 smaller_model.add(layers.Dense(4, activation='relu'))
54 smaller_model.add(layers.Dense(1, activation='sigmoid'))
55 smaller_model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
56
57 smaller_model_hist = smaller_model.fit(x_train, y_train, epochs=20, batch_size=512,
    validation_data=(x_test, y_test))
58
59 smaller_model_val_loss = smaller_model_hist.history['val_loss']
60
61 # b+ is for "blue cross"
62 plt.plot(epochs, original_val_loss, 'b+', label='Original model')
63 # "bo" is for "blue dot"
64 plt.plot(epochs, smaller_model_val_loss, 'bo', label='Smaller model')
65 plt.xlabel('Epochs')
66 plt.ylabel('Validation loss')
67 plt.legend()
68
69 plt.show()
70
```



# Full code (4/5)

```
71 # method 2: adding weight regularization
72 from tensorflow.keras import regularizers
73
74 l2_model = models.Sequential()
75 l2_model.add(layers.Dense(16, kernel_regularizer=regularizers.l2(0.001),
76     activation='relu', input_shape=(10000,)))
77 l2_model.add(layers.Dense(16, kernel_regularizer=regularizers.l2(0.001),
78     activation='relu'))
79 l2_model.add(layers.Dense(1, activation='sigmoid'))
80
81 l2_model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
82 l2_model_hist = l2_model.fit(x_train, y_train, epochs=20, batch_size=512,
83     validation_data=(x_test, y_test))
84
85 l2_model_val_loss = l2_model_hist.history['val_loss']
86
87 plt.plot(epochs, original_val_loss, 'b+', label='Original model')
88 plt.plot(epochs, l2_model_val_loss, 'bo', label='L2-regularized model')
89 plt.xlabel('Epochs')
90 plt.ylabel('Validation loss')
91 plt.legend()
92 plt.show()
```



# Full code (5/5)

```
92 #method 3: adding dropout
93 dpt_model = models.Sequential()
94 dpt_model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
95 dpt_model.add(layers.Dropout(0.5))
96 dpt_model.add(layers.Dense(16, activation='relu'))
97 dpt_model.add(layers.Dropout(0.5))
98 dpt_model.add(layers.Dense(1, activation='sigmoid'))
99
100 dpt_model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
101
102 dpt_model_hist = dpt_model.fit(x_train, y_train, epochs=20, batch_size=512,
    validation_data=(x_test, y_test))
103
104 dpt_model_val_loss = dpt_model_hist.history['val_loss']
105
106 plt.plot(epochs, original_val_loss, 'b+', label='Original model')
107 plt.plot(epochs, dpt_model_val_loss, 'bo', label='Dropout-regularized model')
108 plt.xlabel('Epochs')
109 plt.ylabel('Validation loss')
110 plt.legend()
111
112 plt.show()
```

# Summary

- 3 methods to avoid overfitting:
  - Reduce the network's size
  - Add weight regularization
  - Add dropout

# The universal workflow of machine learning (1/4)

- Defining the problem and assembling a dataset
  - What will your input data be?
  - What are you trying to predict?
  - What type of problem are you facing?
    - Binary classification
    - Multiclass classification
    - Scalar regression
    - Multiclass, multilabel classification
  - Input/output: I/O should be defined
    - What data you' ll have?

# The universal workflow of machine learning (2/4)

- Choosing a measure of success
  - Accuracy
  - Precision/recall
  - ROC: receiver operating characteristic curve
- Deciding on an evaluation protocol
  - Hold-out validation: having plenty of data
  - K-fold cross-validation: too few samples
  - Iterated K-fold validation: highly accurate model, with little data
    - In most cases, the first (hold-out validation) will work well enough.

# The universal workflow of machine learning (3/4)

- Preparing your data
  - The data should be formatted as tensors
  - Values to be scaled to small values
    - $[-1,1]$  or  $[0,1]$
  - Normalizing for different features
  - Feature engineering
- Developing a model that does better than baseline
  - MNIST digits 0~9: baseline accuracy: 0.1; 10%
    - Should be better than random guess

# The universal workflow of machine learning (4/4)

- Last-layer activation, loss function, optimization configuration

**Table 4.1** Choosing the right last-layer activation and loss function for your model

Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	binary_crossentropy
Multiclass, single-label classification	softmax	categorical_crossentropy
Multiclass, multilabel classification	sigmoid	binary_crossentropy
Regression to arbitrary values	None	mse
Regression to values between 0 and 1	sigmoid	mse or binary_crossentropy

- Scaling up: developing a model that overfits
  - Add layers, make the layers bigger, train more epochs
- Regularization: add dropout, add L2 regularization

# Textbook Reading

- *Deep Learning with Python*
  - Ch 4, Fundamentals of machine learning
    - p.94 - p.116